

CN FORESIGHT WORKSHOP

A facilitator's run-deck for the five-method foresight pipeline — turning six analysts' signals into calibrated claims about the future threat surface.

HALF DAY / ~4 HOURS

6 ANALYSTS

OUTPUT → JUNE 16 DECK

AGENDA

01	SCANNING Six analysts' signals → one shared set of named drivers	35 min
02	DRIVERS & UNCERTAINTY Separate the baseline from the two scenario axes	25–30 min
03	SCENARIOS VIA 2x2 Four plausible futures, red-teamed — the centerpiece	75 min
04	TAIL RISK Wild cards + pre-mortem for the blind-spot shifts	30 min
05	INDICATORS & CALIBRATION Convert shifts into monitorable, calibrated claims	30 min +

≈ HALF DAY • ~4 HOURS • OUTPUT FEEDS THE JUNE 16 LEADERSHIP DECK

SCANNING

Turn six analysts' pre-work signals into one shared, de-duplicated set of named drivers — with cross-category links the individual scans could not see.

In: per-analyst STEEP-M scans (3–5 sourced signals)

Out: 8–12 named drivers

Time: 35 min

> RUN IT

- 2 min

Restate the question and the trend-vs-shift distinction. Two rules out loud: every signal needs a source and a one-line "so what."
- 5 min

Silent post — one signal per sticky, under its STEEP-M heading. Atomic, movable units.
- 3 min


Round-robin read-out, 30 seconds each on two strongest signals. No debate yet.
- 8 min

Affinity cluster. Move stickies that share an underlying force together. Push for cross-category clusters: "what force is underneath these?"
- 5 min

Name each cluster as a driver — a force, not an event. Aim for 8–12.
- 2 min

Debrief. Park orphan signals; they are wild-card candidates for Step 4.

STEP 01 · SIGNALS → DRIVERS



The move that matters: a Technology signal and an Attack-surface signal cluster into one driver no single analyst's lens produced. Clustering across categories is where scanning earns its place.

SCANNING

[+] WORKED EXAMPLE

"Agentic AI lowers the cost of tailored phishing" (Technology) and "third-party SaaS concentration rising" (Attack surface) cluster into a single driver: *the cost of credible intrusion is falling while our dependency surface widens.*

[!] WATCH-OUTS

Analysts stay in their lane and miss cross-category links. The "so what" line is where weak signals die — enforce it. Resist writing drivers as events; a driver is a force that shapes many events.

DRIVERS & UNCERTAINTY

Separate the near-certain baseline from the genuinely uncertain forces, and pick the two high-impact, high-uncertainty drivers that become your scenario axes.

In: 8-12 named drivers

Out: baseline trend list + 2 scenario axes

Time: 25-30 min

> RUN IT

- 3 min

Teach the two dimensions. Impact = how much the threat surface changes if the force moves. Uncertainty = how confidently you can predict its direction. They are independent.
- 5 min

Draw the large grid: uncertainty on the horizontal, impact on the vertical.
- 7 min


Place each driver. Hand out dots or place by consensus. Use dot spread to expose genuine disagreement.
- 5 min

Read the grid. Top-left = baseline. Top-right = axis candidates. Bottom half waits.
- 5 min

Select two axes from the top-right: highest impact, genuinely uncertain, and independent of each other.
- 2 min

Sanity check. State each as a spectrum with two plausible poles. If you can predict the winner, it's baseline.

STEP 02 · IMPACT × UNCERTAINTY



Two reads of the grid: *top-left is your "anchored by current priorities" baseline; top-right holds your scenario axes; the bottom waits. Pick axes that are **independent** — or the 2×2 collapses into a line.*

DRIVERS & UNCERTAINTY

[+] WORKED EXAMPLE

"Regulatory disclosure pressure increases" lands high-impact, low-uncertainty → *baseline*. "Does AI shift advantage to attackers or defenders" and "is critical dependency concentrated or distributed" both land high-impact, high-uncertainty, and are independent → *the two axes*.

[!] WATCH-OUTS

Important-but-predictable drivers love to masquerade as uncertainties — demote them to baseline. Correlated axes are the most common failure here, and they wreck Step 3.

SCENARIOS VIA 2x2

Build four plausible, distinct futures from the two axes — each a different threat surface — then stress-test them so they survive a leadership room.

In: 2 independent axes

Out: 4 named scenarios + red-team critique

Time: 75 min

> RUN IT

- 5 min

Draw the 2x2 and label all four poles. Each quadrant is a world where both drivers resolved a particular way.
- 10 min

We-do. Build one quadrant together: name it, sketch what's true, what attacks get cheap, where BlackRock is exposed.
- 30 min


You-do. Pairs take a quadrant: a memorable name, 3–4 sentence narrative, dominant threats, exposure, and the "leading domino."
- 15 min

Gallery walk. Each pair presents its quadrant in three minutes.
- 10 min

Red-team pass. A designated red hat attacks each scenario for consistency, distinctness, and blind spots. Capture every challenge.
- 5 min

Harvest. List shifts that recur across quadrants (robust) vs. unique to one (contingent). Feeds Step 5.

STEP 03 · FOUR FUTURES



The axes are the work. *If your four quadrants reduce to "good vs. bad," the axes were correlated — go back to Step 2. Distinct, surprising quadrants mean you chose well.*

SCENARIOS VIA 2x2

[+] WORKED EXAMPLE

AI advantage × dependency concentration yields four worlds: *Brittle Monoculture* (attacker-favored, concentrated), *Defensible Core* (defender-favored, concentrated), *Diffuse Skirmish* (attacker-favored, distributed), and *Resilient Mesh* (defender-favored, distributed).

[!] WATCH-OUTS

Quadrants collapsing into "good vs. bad" signal correlated axes. Narratives drifting into prediction miss the point — each quadrant is a conditional. And four lightly reskinned versions of today waste the method.

TAIL RISK

Surface the high-impact, low-probability discontinuities the four scenarios missed, and use imagined failure to expose blind-spot shifts. This is where the "be bold" material comes from.

In: 4 scenarios + parked orphan signals

Out: wild-card list + pre-mortem causal chains

Time: 30 min

> RUN IT

WILD CARDS

- 3 min

Teach. Low probability is the entry ticket, not a disqualifier. Pull the parked orphan signals back into view.
- 7 min

Silent brainstorm, then round-robin — one discontinuity each. Ban the word "unlikely."
- 5 min

Plot on a probability × impact grid; keep only the high-impact corner.

PRE-MORTEM


- 2 min

Set the frame: "It's 2028. BlackRock has taken a defining InfoSec hit. It already happened."
- 8 min

Silent write — the headline and the proximate cause. Backward from failure beats forward speculation.
- 5 min

Share and extract. Trace each chain to an early observable — a candidate indicator that maps to a shift.

STEP 04 · WILD CARDS + PRE-MORTEM



Two contrarian passes. Wild cards keep the high-impact corner regardless of odds. The pre-mortem fixes a 2028 failure and reasons backward — the earliest observable in that chain becomes a free indicator for Step 5.

TAIL RISK

[+] WORKED EXAMPLE

A pre-mortem produces: a trusted analytics vendor embedded in the core platform supply chain is compromised and used as a launch point. The backward chain — unreviewed third-party code in a privileged context, anomalous update cadence — exposes a shift: *trust concentrates in a few deeply embedded vendors.*

[!] WATCH-OUTS

The group will try to kill wild cards as too unlikely — that's why you ban the word. Pre-mortems blaming generic "human error" are useless; push for a specific shift. And don't let it become doom theater: every pre-mortem must yield an indicator.

INDICATORS & CALIBRATION


Convert the harvested shifts into monitorable claims with explicit confidence, so what leaves the room is defensible and trackable rather than speculative.

In: harvested shifts (robust, contingent, tail)	Out: 5–7 shifts · 2–3 indicators each · a probability · a horizon
--	--

> RUN IT

8 min Converge. Vote the harvest down to 5–7 shifts. State each as a one-line claim about a change in the threat surface.	8 min Calibrate. Restate each shift with a probability and horizon. Start from an outside-view base rate, then adjust. Record every forecast.
10 min Indicators. "What would we observe in 6–18 months if this is happening?" Write 2–3 concrete, collectible indicators. Assign an owner per shift.	4 min Reveal the sealed pre-work forecasts and compare. Carry live disagreements into the deck as honest analytic-confidence language.

STEP 05 · WATCHABLE + HONEST



Two halves. Indicators make a shift watchable — each must be specific and collectible or it isn't a signpost. Calibration makes a forecast honest; confidence is not calibration, and you only learn the difference by recording forecasts and scoring them.

INDICATORS & CALIBRATION

[+] WORKED EXAMPLE

Shift: AI lowers the cost of tailored intrusion.

Indicators: a rise in high-quality spear-phishing against mid-tier firms; underground listings for AI-assisted tooling; the first breach attributed to agentic tooling. *Claim:* 60% chance that within three years an AI-assisted intrusion is the proximate cause of a material breach at a top-10 asset manager.

[!] WATCH-OUTS

Vague indicators are the most common failure — they must be specific and collectible. Confidence is not calibration; you only learn it by tracking outcomes over time. The forecast record is the first thing dropped under time pressure and the thing with the most evidence behind it — protect it.